

# Optimization

Machine Learning with R

Basel R Bootcamp



October 2019

# Fighting overfitting

When a model **fits the training data too well** on the expense of its performance in prediction, this is called overfitting.

Just because model A is better than model B in training, does not mean it will be better in testing! Extremely flexible models are '**wolves in sheep's clothing**'.

But is there nothing we can do?



adapted from [victoriarollison.com](http://victoriarollison.com)

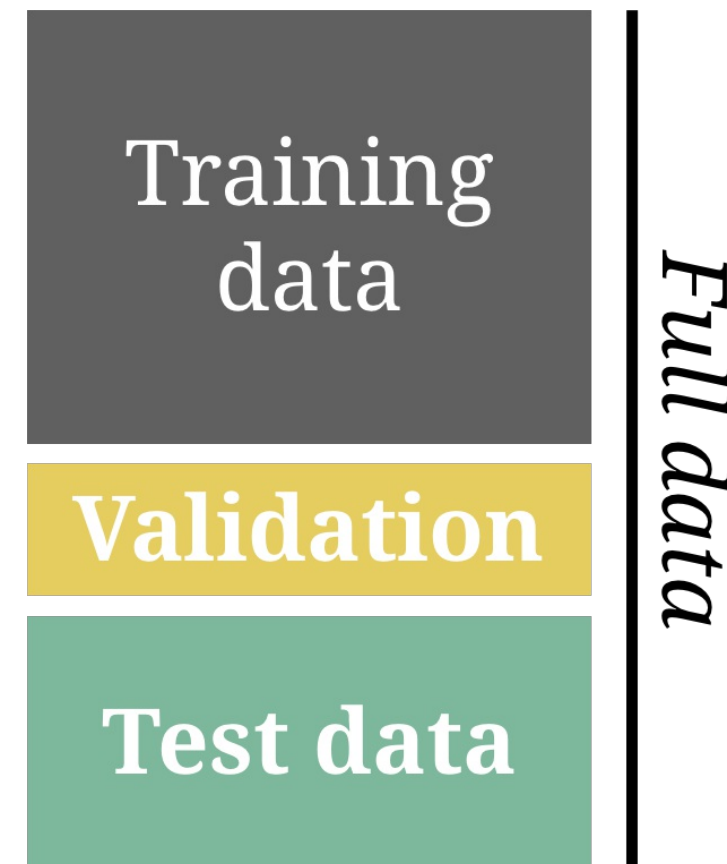
# Tuning parameters

All machine learning models are equipped with tuning parameters that **control model complexity**.

These tuning parameters can be identified using a **validation set** created from the training data.

## Logic

- 1 - Create separate test set.
- 2 - Fit model using various tuning parameters.
- 3 - Select tuning leading to best prediction on validation set.
- 4 - Refit model to entire training set (training + validation).




# Resampling methods


Resampling methods automatize and generalize model tuning.

Method	Description
<i>k-fold cross-validation</i>	Splits the data in $k$ -pieces, use <b>each piece once</b> as the validation set, while using the other one for training.
<i>Bootstrap</i>	For $B$ bootstrap rounds <b>sample</b> from the data <b>with replacement</b> and split the data in training and validation set.

*Cross-validation*  
10-fold



*Bootstrap*

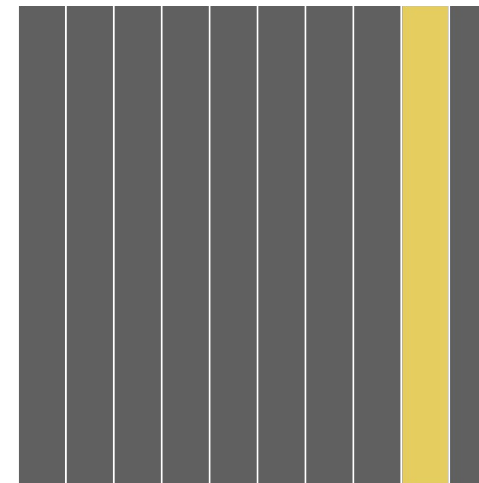


# Resampling methods


Resampling methods automatize and generalize model tuning.

Method	Description
<i>k-fold cross-validation</i>	Splits the data in $k$ -pieces, use <b>each piece once</b> as the validation set, while using the other one for training.
<i>Bootstrap</i>	For $B$ bootstrap rounds <b>sample</b> from the data <b>with replacement</b> and split the data in training and validation set.

*Cross-validation*  
10-fold



*Bootstrap*




# Resampling methods


Resampling methods automatize and generalize model tuning.

Method	Description
<i>k-fold cross-validation</i>	Splits the data in $k$ -pieces, use <b>each piece once</b> as the validation set, while using the other one for training.
<i>Bootstrap</i>	For $B$ bootstrap rounds <b>sample</b> from the data <b>with replacement</b> and split the data in training and validation set.

*Cross-validation*  
10-fold



*Bootstrap*



Regression

Decision Trees

Random Forests

# Regularized regression

Penalizes regression loss for having large  $\beta$  values using the **lambda  $\lambda$  tuning parameter** and one of several penalty functions.

$$\text{Regularized loss} = \sum_i^n (y_i - \hat{y}_i)^2 + \lambda \sum_j^p f(\beta_j)$$

Name	Function	Description
<i>Lasso</i>	$ \beta_j $	Penalize by the <b>absolute</b> regression weights.
<i>Ridge</i>	$\beta_j^2$	Penalize by the <b>squared</b> regression weights.
<i>Elastic net</i>	$ \beta_j  + \beta_j^2$	Penalize by Lasso and Ridge penalties.



from [mallorcazeitung.es](http://mallorcazeitung.es)



# Regularized regression

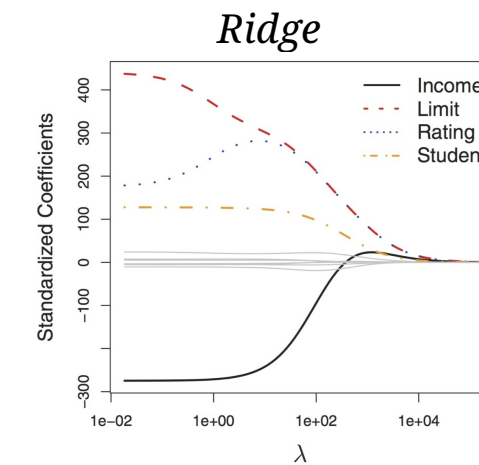
Despite **superficial similarities**, Lasso and Ridge show very different behavior.

## Ridge

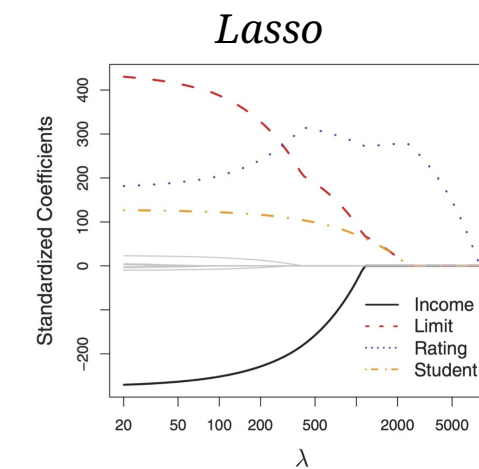
By penalizing the most extreme  $\beta$ s most strongly, Ridge leads to (relatively) more **uniform  $\beta$ s**.

## Lasso

By penalizing all  $\beta$ s equally, irrespective of magnitude, Lasso drives some  $\beta$ s to 0 resulting effectively in **automatic feature selection**.



from James et al. (2013) ISLR



from James et al. (2013) ISLR

# Regularized regression

To fit Lasso or Ridge penalized regression in R, use `method = "glmnet"`.

Specify the **type of penalty** and the **penalty weight** using the `tuneGrid` argument.

tuneGrid settings

Parameter	Description
<code>alpha = 1</code>	Regression with Lasso penalty.
<code>alpha = 0</code>	Regression with Ridge penalty.
<code>lambda</code>	Regularization penalty weight.

```
# Train ridge regression
train(form = criterion ~ .,
      data = data_train,
      method = "glmnet",
      trControl = ctrl,
      tuneGrid =
        expand.grid(alpha = 0,    # Ridge
                    lambda = 1)) # Lambda

# Train lasso regression
train(form = criterion ~ .,
      data = data_train,
      method = "glmnet",
      trControl = ctrl,
      tuneGrid =
        expand.grid(alpha = 1,    # Lasso
                    lambda = 1)) # Lambda
```

Regression

Decision Trees

Random Forests


# Decision trees

Decision trees have a **complexity parameter** called `cp`.

$$\text{Loss} = \text{Impurity} + \text{cp} * (n \text{ terminal nodes})$$

`tuneGrid` settings

Parameter	Description
Small <code>cp</code> , e.g., <code>cp &lt; .01</code>	Low penalty leading to <b>complex trees</b> .
Large <code>cp</code> , e.g., <code>cp &lt; .20</code>	Large penalty leading to <b>simple trees</b> .



# Decision trees

Decision trees have a **complexity parameter** called `cp`.

$$\text{Loss} = \text{Impurity} + \\ \text{cp} * (n \text{ terminal nodes})$$

## tuneGrid settings

Parameter	Description
Small <code>cp</code> , e.g., <code>cp &lt; .01</code>	Low penalty leading to <b>complex trees</b> .
Large <code>cp</code> , e.g., <code>cp &lt; .20</code>	Large penalty leading to <b>simple trees</b> .

```
# Decision tree with a defined cp = .01
train(form = income ~ .,
      data = baselers,
      method = "rpart", # Decision Tree
      trControl = ctrl,
      tuneGrid =
        expand.grid(cp = .01)) # cp

# Decision tree with a defined cp = .2
train(form = income ~ .,
      data = baselers,
      method = "rpart", # Decision Tree
      trControl = ctrl,
      tuneGrid =
        expand.grid(cp = .2)) # cp
```

Regression

Decision Trees

Random Forests


# Random Forest

Random Forests have a **diversity parameter** called `mtry`.

Technically, this controls how many features are randomly considered at each split of the trees.

`tuneGrid` settings

Parameter	Description
Small <code>mtry</code> , e.g., <code>mtry = 1</code>	<b>Diverse forest.</b> In a way, less complex.
Large <code>mtry</code> , e.g., <code>mtry &gt; 5</code>	<b>Similar forest.</b> In a way, more complex.



# Random Forest

Random Forests have a **diversity parameter** called `mtry`.

Technically, this controls how many features are randomly considered at each split of the trees.

tuneGrid settings

Parameter	Description
Small <code>mtry</code> , e.g., <code>mtry = 1</code>	<b>Diverse forest.</b> In a way, less complex.
Large <code>mtry</code> , e.g., <code>mtry &gt; 5</code>	<b>Similar forest.</b> In a way, more complex.


```
# Random forest with a defined mtry = 2
train(form = income ~ .,
      data = baselers,
      method = "rf", # Random forest
      trControl = ctrl,
      tuneGrid =
        expand.grid(mtry = 2)) # mtry

# Random forest with a defined mtry = 5
train(form = income ~ .,
      data = baselers,
      method = "rf", # Random forest
      trControl = ctrl,
      tuneGrid =
        expand.grid(mtry = 5)) # mtry
```



# caret

Parameter tuning with k-fold cross-validation



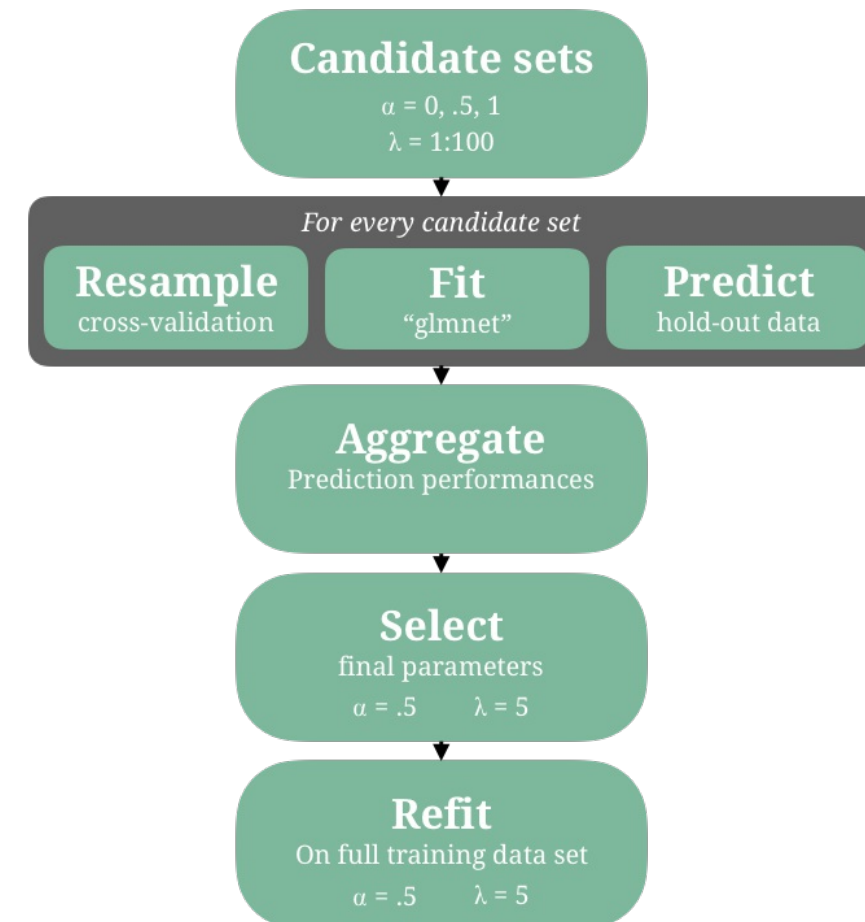
# k-fold cross validation for Ridge and Lasso

## Goal

Use 10-fold cross-validation to identify **optimal regularization parameters** for a regression model.

## Consider

$\alpha \in 0, .5, 1$  and  $\lambda \in 1, 2, \dots, 100$



# trainControl()

Specify the use of k-fold cross-validation using the trainControl() function.

## trainControl() arguments

Argument	Description
----------	-------------

method	The resampling method, use cv for cross validation.
--------	---

number	The number of folds.
--------	----------------------

```
# Specify 10 fold cross-validation
ctrl_cv <- trainControl(method = "cv",
                        number = 10)

# Predict income using glmnet
glmnet_mod <- train(form = income ~ .,
                    data = baselers,
                    method = "glmnet",
                    trControl = ctrl_cv)
```

# tuneGrid

Specify the tuning parameter values to consider using the **tuneGrid**.

tuneGrid expects a **list or data frame** as input.

**Parameter combinations** can be easily created using `expand.grid`.

```
# Specify 10 fold cross-validation
ctrl_cv <- trainControl(method = "cv",
                        number = 10)

# Predict income using glmnet
glmnet_mod <- train(form = income ~ .,
                   data = baselers,
                   method = "glmnet",
                   trControl = ctrl_cv,
                   tuneGrid = expand.grid(
                     alpha = c(0, .5, 1),
                     lambda = 1:100))
```

# k-Fold Cross validation

```
# Print summary information  
glmnet_mod
```

At the end...

RMSE was used to select the optimal model using the smallest value. The final values used for the model were  $\alpha = 1$  and  $\lambda = 27$ .

glmnet

1000 samples  
19 predictor

No pre-processing

Resampling: Cross-Validated (10 fold)

Summary of sample sizes: 900, 901, 900, 901, 901, 899, ...

Resampling results across tuning parameters:

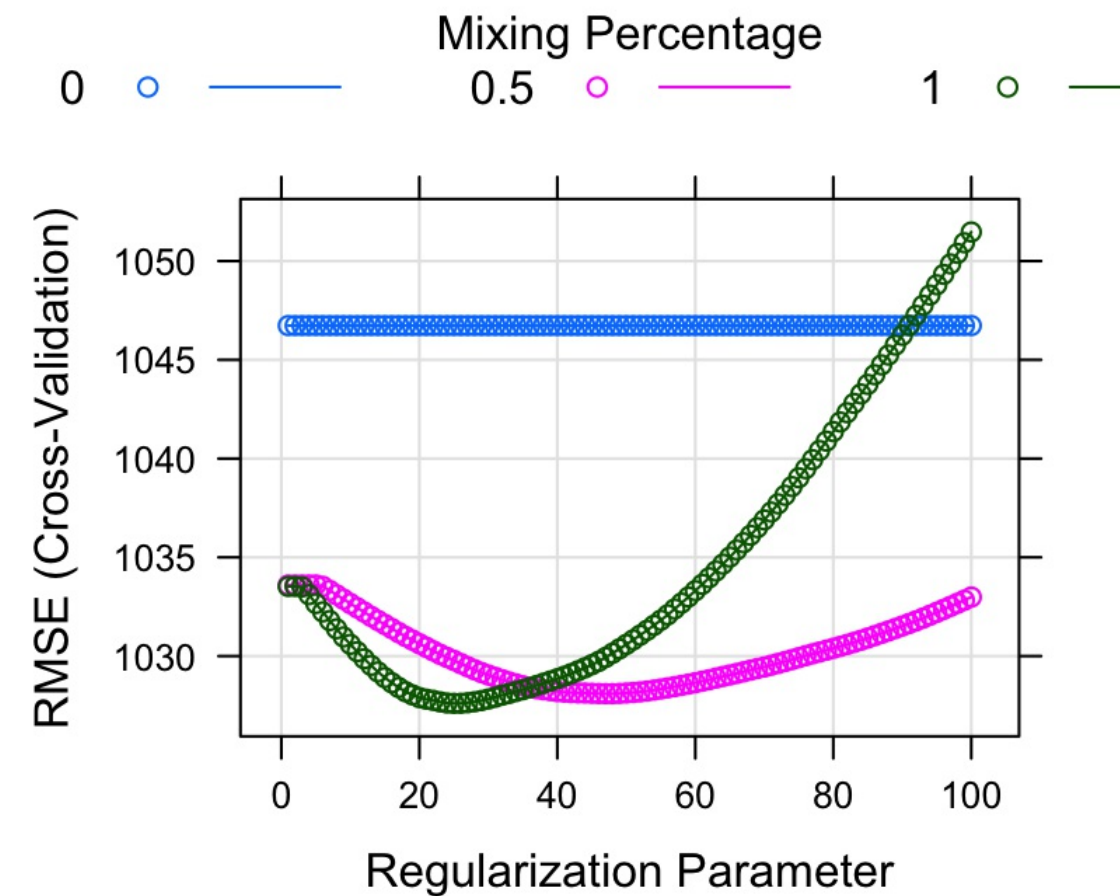
alpha	lambda	RMSE	Rsquared	MAE
0.0	1	1047	0.8614	823.3
0.0	2	1047	0.8614	823.3
0.0	3	1047	0.8614	823.3
0.0	4	1047	0.8614	823.3
0.0	5	1047	0.8614	823.3
0.0	6	1047	0.8614	823.3
0.0	7	1047	0.8614	823.3
0.0	8	1047	0.8614	823.3
0.0	9	1047	0.8614	823.3
0.0	10	1047	0.8614	823.3
0.0	11	1047	0.8614	823.3
0.0	12	1047	0.8614	823.3

# k-Fold Cross validation

```
# Visualise tuning error curve  
plot(glmnet_mod)
```

At the end...

RMSE was used to select the optimal model using the smallest value. The final values used for the model were  $\alpha = 1$  and  $\lambda = 27$ .



# Final model

```
# Model coefficients for best  
# alpha and lambda  
coef(glmnet_mod$finalModel,  
      glmnet_mod$bestTune$lambda)
```

```
25 x 1 sparse Matrix of class "dgCMatrix"  
1  
(Intercept) 462.1958  
id .  
sexmale .  
age 116.1387  
height 1.8865  
weight .  
educationobligatory_school .  
educationSEK_II .  
educationSEK_III 1.1857  
confessionconfessionless 3.9774  
confessionevangelical-reformed .  
confessionmuslim .  
confessionother .  
children -21.0502  
happiness -128.5448  
fitness .  
food 2.3193  
alcohol 22.2351  
tattoos -24.8320  
rhine 0.4256
```

# Model comparison

Compare the prediction performance of several models with `resamples()`.

The `summary()` of this object will print 'prediction' error statistics from cross-validation during training. This is your **estimate of future prediction performance!**

```
# Simple competitor model
glm_mod <- train(form = income ~ .,
                 data = baselers,
                 method = "glm",
                 trControl = ctrl_cv)

# Determine prediction statistics
resamples_mod <- resamples(
  list(glmnet = glmnet_mod,
        glm = glm_mod))

# Print result summary
summary(resamples_mod)
```



# Model comparison

Compare the prediction performance of several models with `resamples()`.

The `summary()` of this object will print 'prediction' error statistics from cross-validation during training. This is your **estimate of future prediction performance!**

```
Call:
summary.resamples(object = resamples_mod)
```

```
Models: glmnet, glm
Number of resamples: 10
```

MAE

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
glmnet	743.1	761.2	818.3	807.8	836.7	891.7	0
glm	734.8	777.7	801.6	812.8	844.5	892.2	0

RMSE

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
glmnet	936.5	990.3	1042	1028	1076	1098	0
glm	950.7	1008.9	1016	1034	1063	1128	0

Rsquared

	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.	NA's
glmnet	0.8386	0.8440	0.8582	0.8638	0.8865	0.9021	0
glm	0.8268	0.8549	0.8694	0.8624	0.8740	0.8825	0

# Practical